


Regional learning: escalation falls as the region learns




Escalation 66% → 2%; local resolution
34% → 82%

Driven by learned knowledge, not tuning
— the routing gate never moved

Local answers resolve in < 50 ms at zero
cloud LLM cost

Confidence dips as the KB widens to
harder events, recovers to 64%

Why the composite moved



LLM reliability and detection evidence saturate at 100%

Retrieval steps up late: Recall@3 54.2%, MRR 50.0%

Routing is weakest at 39.3% — our main improvement target

Severity 71.4%: a complete answer can still misjudge severity